

Learning Deep Representation for Imbalanced Classification

Chen Huang^{1,2} Yining Li¹ Chen Change Loy^{1,3} Xiaoou Tang^{1,3}

¹Department of Information Engineering, The Chinese University of Hong Kong

²SenseTime Group Limited

³Shenzhen Institutes of Advanced Technology, Chinese Academy of Sciences

Fchuang, lly015, cclloy, xtang@ie.cuhk.edu.hk

Abstract

Data in vision domain often exhibit highly-skewed class distribution, i.e., most data belong to a few majority classes, while the minority classes only contain a scarce amount of instances. To mitigate this issue, contemporary classification methods based on deep convolutional neural network (CNN) typically follow classic strategies such as class re-sampling or cost-sensitive training. In this paper, we conduct extensive and systematic experiments to validate the effectiveness of these classic schemes for representation learning on class-imbalanced data. We further demonstrate that more discriminative deep representation can be learned by enforcing a deep network to maintain both inter-cluster and inter-class margins. This tighter constraint effectively reduces the class imbalance inherent in the local data neighborhood. We show that the margins can be easily deployed in standard deep learning framework through quintuplet instance sampling and the associated triple-header hinge loss. The representation learned by our approach, when combined with a simple k-nearest neighbor (kNN) algorithm, shows significant improvements over existing methods on both high- and low-level vision classification tasks that exhibit imbalanced class distribution.

1. Introduction

Many data in computer vision domain naturally exhibit imbalance in their class distribution. For instance, the number of positive and negative face pairs in face verification is highly skewed since it is easier to obtain face images of different identities (negative) than faces with matched identity (positive) during data collection. In face attribute recognition [25], it is comparatively easier to find persons with “normal-sized nose” attribute on web images than that of “big-nose”. For image edge detection, the image edge structures intrinsically follow a power-law distribution, e.g., horizontal and vertical edges outnumber those with “Y” shape. Without handling the imbalance issue conventional meth-

ods tend to be biased toward the majority class with poor accuracy for the minority class [18].

Deep representation learning has recently achieved great success due to its high learning capacity, but still cannot escape from such negative impact of imbalanced data. To counter the negative effects, one often chooses from a few available options, which have been extensively studied in the past [7, 9, 11, 17, 18, 30, 40, 41, 46, 48]. The first option is re-sampling, which aims to balance the class priors by under-sampling the majority class or over-sampling the minority class (or both). For instance, Oquab et al. [32] resample the number of foreground and background image patches for learning a convolutional neural network (CNN) for object classification. The second option is cost-sensitive learning, which assigns higher misclassification costs to the minority class than to the majority. In image edge detection, for example, Shen et al. [35] regularize the softmax loss to cope with the imbalanced edge class distribution. Other alternatives exist, e.g., learning rate adaptation [18].

Are these methods the most effective way to deal with imbalance data in the context of deep representation learning? The aforementioned options are well studied for the so called ‘shallow’ model [12] but their implications have not yet been systematically studied for deep representation learning. Importantly, such schemes are well-known for some inherent limitations. For instance, over-sampling can easily introduce undesirable noise with overfitting risks, and under-sampling is often preferred [11] but may remove valuable information. Such nuisance factors can be equally applicable to deep representation learning.

In this paper, we wish to investigate a better approach for learning a deep representation given class-imbalanced data. Our method is motivated by the observation that the minority class often contains very few instances with high degree of visual variability. The scarcity and high variability make the genuine neighborhood of these instances easy to be invaded by other imposter nearest neighbors¹. To this end,

¹An imposter neighbor of a data point x_i is another data point x_j with a different class label, $y_i \neq y_j$.

we propose to learn an embedding $f(x) \in \mathbb{R}^d$ with a CNN to ameliorate such invasion. The CNN is trained with instances selected through a new quintuplet sampling scheme and the associated triple-header hinge loss. The learned embedding produces features that preserve not only locality across the same-class clusters but also discrimination between classes. We demonstrate that such “quintuplet loss” introduces a tighter constraint for reducing imbalance in the local data neighborhood when compared to existing triplet loss. We also study the effectiveness of classic schemes of class re-sampling and cost-sensitive learning in our context.

Our key contributions are as follows: (1) we show how to learn deep feature embeddings for imbalanced data classification, which is understudied in the literature; (2) we formulate a new quintuplet sampling method with the associated triple-header loss that preserves locality across clusters and discrimination between classes. Using the learned features, we show that classification can be simply achieved by a fast cluster-wise kNN search followed by a local large margin decision. The proposed method, called Large Margin Local Embedding (LMLE)-kNN, achieves state-of-the-art results in the large-scale imbalanced classification tasks of (binary) face attributes and (multi-class) image edges.

2. Related Work

Previous efforts to tackle the class imbalance problem can be mainly divided into two groups: data re-sampling [7, 11, 17, 18, 30] and cost-sensitive learning [9, 40, 41, 46, 48]. The former group aims to alter the training data distribution to learn equally good classifiers for the majority and minority classes, usually by random under-sampling and over-sampling techniques. The latter group, instead of manipulating samples at the data level, operates at the algorithmic level by adjusting misclassification costs. A comprehensive literature survey can be found in [18].

A well-known issue with replication-based random over-sampling is its tendency to overfit. More radically, it does not actually increase any information, and fails in solving the fundamental “lack of data” problem. To address this, SMOTE [7] creates new non-replicated examples by interpolating neighboring minority class instances. Several variants of SMOTE [17, 30] followed for improvements. However, their broadened decision regions are still error-prone by synthesizing noisy and borderline examples. Therefore under-sampling is often preferred to over-sampling [11], although potentially valuable information may be removed. Cost-sensitive alternatives avoid these problems by directly imposing heavier penalty on misclassifying the minority class. For example, in [40] the classic SVM is made cost-sensitive to improve classification on highly skewed datasets. Zadrozny et al. [46] combined cost sensitivity with ensemble approaches to further improve classification accuracy. Many other methods follow this practice of designing

classifier ensemble to combat imbalance (e.g., [9, 41]), and boosting [41] offers an easy way to embed the costs by updating example weights. Chen et al. [9] resorted to bagging which is less vulnerable to noise than boosting, and generated a cost-sensitive version of random forest.

None of the above works addresses the class imbalance learning using CNN. They rely on shallow models and hand-crafted features. To our knowledge, only few works [21, 22, 48] approach imbalanced classification via deep learning. Jeatrakul et al. [21] treated the Complementary Neural Network as an under-sampling technique, and combined it with SMOTE over-sampling to balance training data. Zhou and Liu [48] studied data resampling in training cost-sensitive neural networks. Khan et al. [22] further seek for joint optimization of the class-sensitive costs and deep features. These works can be seen as natural extensions to existing imbalanced learning techniques, while neglecting the underlying data structure for discriminating imbalanced data. Motivated from this, we propose a “data structure-aware” deep learning approach with built-in margins for imbalanced classification, where the classic schemes of data resampling and cost-sensitive learning are also studied systematically.

Attribute recognition: Face attributes are useful as mid-level features for many applications like face verification [3, 24, 25]. It is challenging to predict them from unconstrained face images due to the large facial variations such as pose and lighting. Most existing methods for attribute recognition extract hand-crafted features from images, which are then fed into some classifiers to predict the presence of an array of face attributes, e.g., “male”, “smile”, etc. Examples are [24, 25] where HOG-like features are extracted on various local face regions to predict attributes. Recent deep learning methods [29, 47] excel by learning powerful features. Zhang et al. [47], for instance, trained pose-normalized CNNs for deep attribute modeling. These studies, however, share a common drawback: they neglect the class imbalance issue in those relatively rare attributes like “big nose” and “bald”. Thus skewed results are expected when predicting the positive class of these under-represented attributes.

Edge detection: State-of-the-art edge detection methods [1, 2, 10, 16, 20, 27, 33] mostly use engineered gradient features to classify edge patches against non-edges. Due to the large variety of edge structures, such a binary classification problem is usually transformed to a multi-class one. [27, 35] first cluster edge patches into hundreds of subclasses, then the goal becomes predicting whether an input patch belongs to each edge subclass or the non-edge class. The final binary task can be accomplished by ensembling classification scores, which works well under the condition of equal amounts of edge and non-edge samples used. Built on the same assumption, recent CNN-based methods

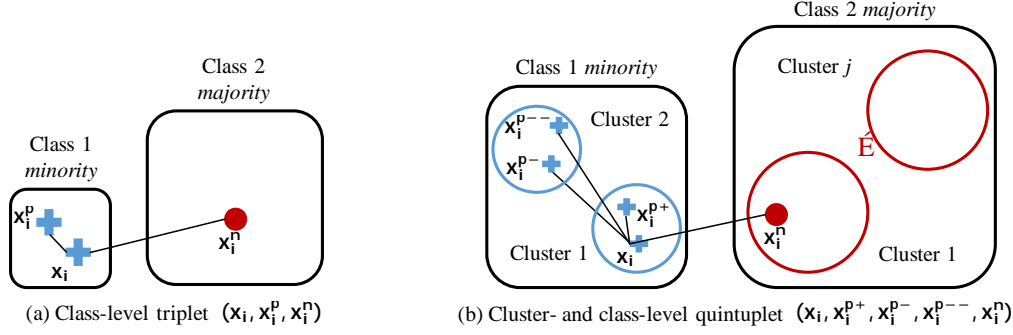


Figure 1. Embeddings by (a) triplet vs. by (b) quintuplet. We exemplify the class imbalance by two different sized classes, where the clusters are formed by k-means. Our quintuplets enforce both inter-cluster and inter-class margins, while triplets only enforce inter-class margins irrespective of different class sizes and variations. This difference leads the unique capability of quintuplets in preserving discrimination in any local neighborhood, and forming a local classification boundary that is insensitive to imbalanced class sizes.

[4, 5, 13, 23, 35, 45] achieve better results by learning deep features. However, all these methods would face the problem of data imbalance between each edge subclass and the dominant non-edge class, which is barely addressed properly. Only Shen et al. [35] attempt to regularize the multi-way softmax loss with a balanced weighting between the “positive” super-class and “negative” class, which is a compromise between the binary and multi-class tasks. Here we propose to explicitly learn discriminative features from imbalanced data.

3. Learning Deep Representation from Class-Imbalanced Data

Given an imagery dataset with imbalanced class distribution, our goal is to learn a Euclidean embedding $f(x)$ from an image x into a feature space R^d , such that the embedded features are discriminative without any possible local class imbalance. We constrain this embedding to live on a d -dimensional hypersphere, i.e., $\|f(x)\|_2 = 1$.

3.1. Quintuplet Sampling

To achieve the aforementioned goal, we select quintuplets from the imbalanced data as illustrated in Fig. 1. Each quintuplet is defined as:

- x_i : an anchor,
- x_i^{p+} : the anchor’s most distant within-cluster neighbor,
- x_i^{p-} : the nearest within-class neighbor of the anchor, but from a different cluster,
- x_i^{p--} : the most distant within-class neighbor of the anchor,
- x_i^n : the nearest between-class neighbor of the anchor.

We wish to ensure that the following relationship holds in the embedding space:

$$\begin{aligned} D(f(x_i), f(x_i^{p+})) &< D(f(x_i), f(x_i^{p-})) \\ &< D(f(x_i), f(x_i^{p--})) < D(f(x_i), f(x_i^n)), \end{aligned} \quad (1)$$

where $D(f(x_i), f(x_j)) = \|f(x_i) - f(x_j)\|_2^2$ is the Euclidean distance.

Such a fine-grained similarity defined by quintuplets has two merits: 1) The ordering in Eq. (1) provides richer information and a stronger constraint than the conventional class-level image similarity. In the latter, two images are considered similar as long as they belong to the same category. In contrast, we require two instances to be close at both class- and cluster-levels to be considered similar. This actually helps build a local classification boundary with the most discriminative local samples. Other irrelevant samples in a class are effectively “ignored” for class separation, making the local boundary robust and insensitive to imbalanced class sizes. 2) The quintuplet sampling is repeated during CNN training, thus avoiding large information loss as in traditional random under-sampling. When compared with over-sampling strategies, it introduces no artificial noise. In practice, to ensure adequate learning for all classes, we collect quintuplets for equal numbers of minority- and majority-class samples x_i in one mini-batch. Section 5 will quantify the efficacy of this re-sampling scheme.

Note in the above, we implicitly assume the imbalanced data are already clustered so that quintuplets can be sampled. In practice, we obtain the initial clusters for each class by applying k-means on some prior features (e.g., for face attribute recognition, we employ the pre-trained DeepID2 features [39] on the face verification task). To make the clustering more robust, an alternating scheme is formulated to refine the clusters using features extracted from the proposed model itself every 5000 iterations. The overall pipeline will be summarized in Section 3.2.

3.2. Triple-Header Hinge Loss

To enforce the relationship in Eq. 1 during feature learning, we apply the large margin idea using the sampled quintuplets. A triple-header hinge loss is formulated to constrain three margins between the four distances, and we de-

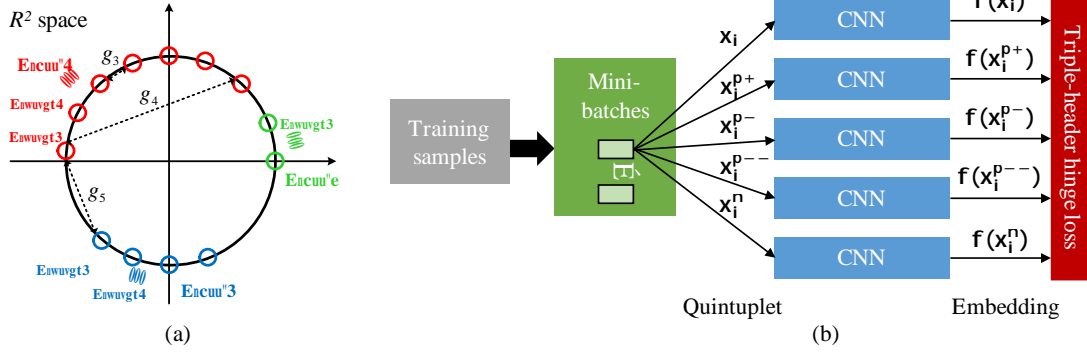


Figure 2. (a) Feature distribution in 2D space and the geometric intuition of margins. (b) Our learning network architecture.

fine the following objective function with slack allowed:

$$\begin{aligned} \min \quad & \sum_i (s_i + t_i + u_i) + \frac{W}{2} \sum_i s_i^2, \\ \text{s.t. :} \quad & \max_i 0, g_1 + D(f(x_i), f(x_i^{p+})) - D(f(x_i), f(x_i^{p-})) - s_i, \\ & \max_i 0, g_2 + D(f(x_i), f(x_i^{p-})) - D(f(x_i), f(x_i^{p--})) - t_i, \\ & \max_i 0, g_3 + D(f(x_i), f(x_i^{p--})) - D(f(x_i), f(x_i^n)) - u_i, \\ & s_i, t_i, u_i \geq 0, \quad s_i \leq 0, \quad t_i \leq 0, \quad u_i \leq 0 \end{aligned} \quad (2)$$

where s_i, t_i, u_i are the slack variables, g_1, g_2, g_3 are the margins, W represents the parameters of the CNN embedding function $f(\cdot)$, and $\frac{W}{2}$ is a regularization parameter.

This formulation effectively regularizes the deep representation learning based on the ordering specified in quintuplets, imposing a tighter constraint than triplets [8, 34, 42]. Ideally, in the hypersphere embedding space, clusters should collapse into small neighborhoods with safe margin g_1 between one another, and g_2 being the largest within a class, and their containing class is also well separated by a large margin g_3 from other classes (see Fig. 2(a)). A merit of our learning algorithm is that the margins can be explicitly determined by a geometric intuition. Suppose there are L training samples in total, class c is of size $L_c, c = 1, \dots, C$. Let all the classes constitute $s \in [0, 1]$ of the entire hypersphere, and we generate clusters of equal-size l for each class. Obviously, the margins' lower bounds are zero. For their upper bounds, g_1^{\max} is obtained when all L/l clusters are squeezed into single points on a proportion s of the sphere. Hence $g_1^{\max} = 2 \sin(\pi s l / L)$, and g_2^{\max} can be approximated as $2 \sin(\pi s (L_c - l) / L)$ via triangle inequality. $g_3^{\max} = 2 \sin(\pi / C)$ when all classes collapse into single points. In practice, we try several decreasing margins by a coarse grid search before actual training.

The learning network architecture is shown in Fig. 2(b). Given a class evenly re-sampled mini-batch, we retrieve for each x_i in it a quintuplet by using a lookup table computed offline. To generate a table of meaningful and discriminative quintuplets, instead of selecting the “hardest” ones from the entire training set, we select “semi-hard” ones by

computing distances on a random subset (50%) of training data to avoid those mislabelled or poor quality data. Then each quintuplet member is fed independently into five identical CNNs with shared parameters. Finally, the output feature embeddings are L_2 normalized and used to compute a triple-header hinge loss by Eq. 2. Back-propagation is used to update the CNN parameters.

To further ensure equal learning for the imbalanced classes, we assign samples in each mini-batch costs such that the class weights therein are identical. The specific cost-sensitive schemes for different classification tasks will be detailed in Section 5 with supporting experiments. Below we summarize the learning steps of the proposed LMLE approach:

1. Cluster for each class by k-means using the learned features from previous round of alternation. For the first round, we use hand-crafted features or prior features obtained from other pre-trained network.
2. Generate a quintuplet table using the cluster and class labels from a random subset (50%) of training data.
3. For CNN training, repeatedly sample mini-batches equally from each class and retrieve the corresponding quintuplets from the offline table.
4. Feed all quintuplets into five identical CNNs to compute the loss in Eq. 2 with cost-sensitivities.
5. Back-propagate the gradients to update the CNN parameters and feature embeddings.
6. Alternate between steps 1-2 and 3-5 every 5000 iterations until convergence (empirically within 4 rounds when we observe no improvements on validation set).

3.3 Differences between “Quintuplet” Loss and Triplet Loss

The triplet loss is inspired by Dimensionality Reduction by Learning an Invariant Mapping (DrLIM) [15] and Large Margin Nearest Neighbor (LMNN) [44]. It is widely used in many recent vision studies [8, 34, 42], aiming to bring data of the same class closer, while data of different classes further away (see Fig. 1). To enforce such a relationship,




Figure 3. From left to right: 2D feature embedding of one imbalanced binary face attribute using DeepID2 model [39] (i.e., prior feature model), triplet-based embedding, quintuplet-based LMLE. We only show 2 Positive Clusters (PC) and 5 Negative Clusters (NC) out of a total of 499 clusters to represent the imbalance.

one needs to generate mini-batches of triplets, i.e., an anchor x_i , a positive instance x_i^p of the same class, and a negative instance x_i^n of different class, for deep feature learning. We argue that they are rather limited in capturing the embedding structure of imbalanced data. Specifically, the similarity information is only extracted at the class-level, which would homogeneously collapse each class irrespective of their different degrees of variation. As a result, the class structures are lost. When a class has high data variability, it is also hard to maintain the class-wise margin, leading to potential invasion of imposter neighbors or even domination of the majority class in local neighborhood. By contrast, the proposed LMLE generates diverse quintuplets that differ in the membership of both clusters and classes. It thus captures the considerable data variability within each class, and can easily enforce the local margin to reduce any local class imbalance. It is worth noting that Wang et al. [42] also aim to learn fine-grained similarity within classes but they do not explicitly model the within-class variations by clustering like us. Fig. 3 illustrates our advantage. Section 5 will further quantify the benefits of our quintuplet loss.

4. Nearest Neighbor Imbalanced Classification

The above LMLE approach offers crucial feature representations for the following classification to perform well on imbalanced data. We choose the simple kNN classifier to show the efficacy of our learned features. Better performance is expected with the use of more elaborated classifiers.

A traditional kNN classifier predicts the class label of query q as the majority label among its kNN in the training set $P = \{(x_i, y_i)\}_{i=1}^L$, where $y_i = 1, \dots, C$ is the (binary or multi-way) class label of sample x_i . Such kNN rule is appealing due to its non-parametric nature, and it is easy to extend to new classes without retraining. However, the underlying equal-class-density assumption is not satisfied and will greatly degrade performance in our imbalanced case.

Hence we modify the kNN classifier in two ways: 1) In the well-clustered LMLE space learned in training stage, we treat each cluster as a single class-specific exemplar²,

²Clustering to aid classification is common in the literature [6, 43].

and perform a fast cluster-wise kNN search. 2) Let (q) be query q 's local neighborhood defined by its kNN cluster centroids $\{m_i\}_{i=1}^k$. We seek a large margin local boundary among (q) , labelling q as the class to which the maximum cluster distance is smaller than the minimum cluster distance to any other class by the largest margin:

$$y_q = \arg \max_{c=1, \dots, C} \min_{\substack{m_j \in (q) \\ y_j = c}} D(f(q), f(m_j)) - \max_{\substack{m_i \in (q) \\ y_i \neq c}} D(f(q), f(m_i)) \quad (3)$$

This large margin local decision offers two advantages:

(i) More resistance to data imbalance: Recall that we fix the cluster size l rather than the number of clusters for each class (to avoid large quantization errors for the minority classes). Thus all the L_c/l clusters from different classes $c = 1, \dots, C$ still exhibit class imbalance. But the large margin rule in Eq. 3 can well solve this issue because it is independent of the cluster number in each class. It is also very suited to our LMLE representation which is learned under the same large margin rule.

(ii) Fast speed: Decision by cluster-wise kNN search is much faster than by sample-wise search.

Finally we summarize the steps for our kNN-based imbalanced classification: for query q ,

1. Find its kNN cluster centroids $\{m_i\}_{i=1}^k$ from all classes learned in the training stage.
2. If all the k cluster neighbors belong to the same class, q is labelled by that class and exit.
3. Otherwise, label q as y_q using Eq. 3.

Note the cluster-wise kNN search can be further sped up using some conventional tricks. For instance, we utilize the KD-tree [36] whose runtime is logarithmic in the number of all clusters (L/l) with a complexity of $O(L/l \log(L/l))$ in comparison to $O(L \log L)$ by sample-wise search. This leads to up to three orders of magnitude speedup over standard kNN classification in practice, making it easy to scale to large datasets with a large number of classes.

4.1. Discussion

There are many previous studies of identifying nearest neighbors that focus on distance metric learning [8, 14, 44], but barely on imbalanced learning. Distance metric learning improves kNN search by optimizing parametric distance functions. For instance, Globerson and Roweis [14] learn the Mahalanobis distance which requires all the same-class samples collapse to one point. However, those classes with high data variability cannot be effectively squeezed to single points [31]. In comparison, our LMLE-kNN captures

within-class variations across clusters, only which are “collapsed” in the embedding space, thus offering both accuracy and speed advantages for the kNN search. In [44], the Mahalanobis distance is learned to directly improve the large margin kNN classification, but not under the imbalanced circumstances. Liu and Chawla [28] provide a weighting scheme over kNN to correct the classification bias. We “correct” this bias more effectively by a large margin kNN classifier with accordingly learned deep features.

5. Results

We study the high-level face attribute classification task and low-level edge detection task, both with large-scale imbalanced data. The originally balanced MNIST digit classification is also studied to experiment with controlled class imbalance. Our face attribute is binary, with severely imbalanced positive and negative samples (e.g., “Bald” attribute: 2% vs. 98%). Our approach predicts 40 attributes simultaneously in a multi-task framework. Edge detection is cast as a multi-class classification problem to address the diversity of edges, i.e., to predict whether an input image patch belongs to any edge class (shape) or the non-edge class. Since the ultimate goal is still binary, the “positive” edge patches and “negative” non-edge patches are usually equally sampled for training. Thus severe imbalance exists between the edge classes (power-law) and dominant non-edge class.

Datasets and evaluation metrics: For face attributes, we use CelebA dataset [29] that contains 10000 identities, each with about 20 images. Every face image is annotated with 40 attributes and 5 key points to align it to 55×47 pixels. We partition the dataset following [29]: the first 160 thousand images (i.e., 8000 identities) for training (10 thousand images for validation), the following 20 thousand for training SVM classifiers for the PANDA [47] and ANet [29] methods, and remaining 20 thousand for testing. To account for the imbalanced positive and negative attribute samples, a balanced accuracy is adopted, that is $\text{accuracy} = 0.5(t_p/N_p + t_n/N_n)$, where N_p and N_n are the numbers of positive and negative samples, while t_p and t_n are the numbers of true positive and true negative. Note that this evaluation metric differs from that employed in [29], where $\text{accuracy} = ((t_p + t_n)/(N_p + N_n))$, which can be biased to the majority class.

For edge detection, we use the BSDS500 [1] dataset that contains 200 training, 100 validation and 200 testing images. We sample 2 million 45×45 pixel training patches, where the numbers of edge and non-edge ones are the same and there are 150 edge classes formed by k-means. Note the ultimate goal of edge detection is to produce a full-scale edge map for one input image, instead of predicting edge classes for local patches by Eq. 3. To estimate the edge map in a robust way, we follow [13] to simply transfer and fuse the overlapping edge label patches of the nearest neighbors

Table 1. Implementation details for the imbalanced classification tasks considered. For each task, we list its used CNN architecture (top left), prior features for clustering (top right), and class-specific cost (bottom).

Face attributes	same w.r.t. [39]	DeepID2 features in [39]
	inv. to class size in batches	
Image edges	same w.r.t. [35]	low-level features in [10]
	inv. to class size in batches & positive sharing weight [35]	
MNIST digits	same w.r.t. [38]	pre-trained features by softmax
	inv. to class size in batches	

found using our learned representations. After applying non-maximal suppression, we evaluate edge detection accuracy by: fixed contour threshold (ODS), per-image best threshold (OIS), and average precision (AP) [1].

The MNIST experiments are carried out on the challenging dataset of MNIST-rot-back-image [26]. This extension dataset contains 28×28 digit images with large rotations and random backgrounds. There are 12000 training images and 50000 testing images, and we augment the training set 10 times by randomly rotating, mirroring and resizing images, leaving 10% in it for validation. We report the mean per-class accuracy in the artificial imbalanced settings.

Parameters: Our CNN is trained using batch size 40, momentum 0.9, and $\eta = 0.0005$ in Eq. 2. We form clusters for each class with sizes around $l = 200$. For those homogeneous classes (with only one possible cluster), our within-class margins actually become nearly zero, which hurts no performance but only reduces quintuplets to the triplets as a lower bound baseline. We search $k = 20$ nearest clusters (i.e., $|q| = 20$ in Eq. 3) for querying. Other task-specific settings are summarized in Table 1. Note the prior features for clustering are not critical to the final results because we will gradually learn deep features in alternation with their clustering every 5000 iterations. Different prior features generally converge to similar results, but at different speeds.

5.1. Comparison with State-of-the-Art Methods

Table 2 compares our LMLE-kNN method for multi-attribute classification with the state-of-the-art Triplet-kNN [34], PANDA [47] and ANet [29] methods, which are trained using the same images and tuned to their best performance. The attributes and their mean per-class accuracies are given in the order of ascending class imbalance level ($= |\text{positive class rate} - 50\%|$) to reflect its impact on performance. It is shown that LMLE-kNN consistently outperforms other methods across all face attributes, with an average gap of 4% over the runner-up ANet. Considering most face attributes exhibit high class imbalance with an average positive class rate of only 23%, such improvements are nontrivial and prove our features’ representation power on imbalanced data. Although the competitive PANDA and

Table 2. Mean per-class accuracy (%) and class imbalance level (= 1/positive class rate-50%) of each of the 40 face attributes on CelebA dataset [29]. Attributes are sorted in an ascending order by the imbalance level. To account for the imbalanced positive and negative attribute samples, a balanced accuracy is adopted, unlike [29]. The results of ANet are therefore different from that reported in [29].

	Attractive	Mouth Open	Smiling	Wear Lipstick	High Cheekbones	Male	Heavy Makeup	Wavy Hair	Oval Face	Pointy Nose	Arched Eyebrows	Black Hair	Big Lips	Big Nose	Young	Straight Hair	Brown Hair	Bags Under Eyes	Wear Earrings	No Beard	Bangs
Imbalance level	1	2	2	3	5	8	11	18	22	22	23	26	26	27	28	29	30	30	31	33	35
Triplet-kNN [34]	83	92	92	91	86	91	88	77	61	61	73	82	55	68	75	63	76	63	69	82	81
PANDA [47]	85	93	98	97	89	99	95	78	66	67	77	84	56	72	78	66	85	67	77	87	92
ANet [29]	87	96	97	95	89	99	96	81	67	69	76	90	57	78	84	69	83	70	83	93	90
LMLE-kNN	88	96	99	99	92	99	98	83	68	72	79	92	60	80	87	73	87	73	83	96	98
	Blond Hair	Bushy Eyebrows	Wear Necklace	Narrow Eyes	5 o'clock Shadow	Receding Hairline	Wear Necktie	Eyeglasses	Rosy Cheeks	Goatee	Chubby	Sideburns	Blurry	Wear Hat	Double Chin	Pale Skin	Gray Hair	Mustache	Bald		Average
Imbalance level	35	36	38	38	39	42	43	44	44	44	44	44	45	45	45	46	46	46	48		
Triplet-kNN [34]	81	68	50	47	66	60	73	82	64	73	64	71	43	84	60	63	72	57	75		72
PANDA [47]	91	74	51	51	76	67	85	88	68	84	65	81	50	90	64	69	79	63	74		77
ANet [29]	90	82	59	57	81	70	79	95	76	86	70	79	56	90	68	77	85	61	73		80
LMLE-kNN	99	82	59	59	82	76	90	98	78	95	79	88	59	99	74	80	91	73	90		84

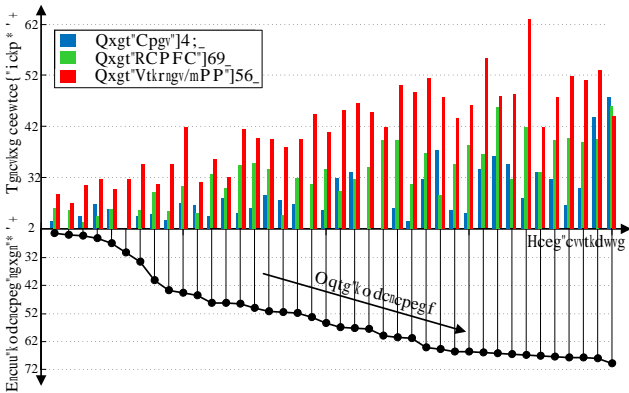


Figure 4. Relative accuracy gains over competitors on the sorted 40 face attributes in Table 2.

ANet are capable of learning a robust representation by ensembling and multi-task learning respectively, they ignore the imbalance issue and thus struggle for highly imbalanced attributes, e.g., “Bald”. Compared with the closely related triplet sampling method [34], our quintuplet sampling better preserves the embedding discrimination on imbalanced data. The advantage is more evident while observing the relative accuracy gains over other methods in Fig. 4. The gains tend to increase with a higher class imbalance level.

Table 3 reports our edge detection results under three metrics. Compared with the classification-based Sketch Token [27] and DeepContour [35] methods, our LMLE-kNN outperforms by a large margin due to the explicit handling of class imbalance during feature learning. Our ap-

Table 3. Edge detection results on BSDS500 [1]. In the bottom cell are deep learning-based methods. *ICCV paper results.

Method	ODS	OIS	AP
gPb-owt-ucm [1]	0.73	0.76	0.73
Sketch Token [27]	0.73	0.75	0.78
SCG [33]	0.74	0.76	0.77
PMI+sPb [20]	0.74	0.77	0.78
SE [10]	0.75	0.77	0.80
OEF [16]	0.75	0.77	0.82
SE+multi-ucm [2]	0.75	0.78	0.76
DeepNet [23]	0.74	0.76	0.76
N ⁴ -Fields [13]	0.75	0.77	0.78
DeepEdge [4]	0.75	0.77	0.81
DeepContour [35]	0.76	0.77	0.80
HFL [5]	0.77	0.79	0.80
CS-SE+DSNA [19]	0.77	0.79	0.81
HED [45]*	0.78	0.80	0.83
LMLE-kNN	0.78	0.79	0.83

proach provides more robust feature representations than only weighting in [35]. This is partially validated by the fact that while [27] and [35] need to train random forests on features to achieve competitive results, ours only relies on kNN label transfer as in [13], but with better results. Fig. 5 shows their visual differences, where LMLE-kNN can accurately discover fine rare edges as well as the majority non-edges that make edge maps clean. Sketch Token [27] and DeepContour [35] suffer from the prediction bias with noisy edges and relatively low recall of fine edges respectively. LMLE-kNN also outperforms [13] and our previous work in [19] by learning effective deep features from imbalanced data. Recent HFL [5], HED [45] methods utilize large VG-




Figure 5. Edge detection examples on BSDS500 [1]. From left to right: input image, ground truth, Sketch Token [27], DeepContour [35], LMLE-kNN. Note the visual differences in the red box.

GNet [37] (13-16 layers) for good performance. In contrast, we only use a lightweight network as in [35] (6 layers).

Computation time: LMLE-kNN takes about 4 days to train for 4 alternative rounds on GPU, and 10ms per sample to extract features. The cluster-wise kNN search is typically 1000 times faster than standard kNN. This enables real-time application to large-scale problems, e.g. the above two with hundreds of thousands to millions of samples.

5.2. Ablation Tests and Control Experiments

Table 4 quantifies the benefits of our quintuplet loss and the re-sampling and cost-sensitive schemes. On both imbalanced tasks, we find favorable performance using the classic schemes, while the proposed quintuplet loss leads to much larger performance gains over baselines. This strongly supports the necessity of imposing additional cluster-wise relations in our quintuplets. Such constraints can better preserve local class structures than triplets, which is critical for ameliorating the invasion of imposter neighbors. Note when the cost-sensitive scheme is applied to strictly balanced re-sampled mini-batches, it would have no effects since the class weights are already equal. However in the case of predicting multiple face attributes, class-balanced data for one attribute will be almost certainly imbalanced for the other attributes, whose class costs can help then. In the case of edge detection, costs can help by further regularizing the impacts of the positive super-class and negative class.

To highlight the effects of our large margin cluster-wise kNN classifier, we replace it with the regular instance-wise kNN classifier in our full method ‘Quintuplet+resample+cost’. As expected, we observe much lower speed and worse results (81% for attributes, 0.77 for edge).

We finally conduct control experiments on the MNIST-rot-back-image dataset [26]. It is originally balanced among 10 digit classes, and we form a Gaussian-like imbalanced class distribution by randomly removing data with increasing amounts (thus more imbalanced). We compare the mean per-class accuracies between 3 baselines in Table 5, where

Table 4. Ablation tests on attributes classification (average accuracy - %) and edge detection (ODS).

Methods	Softmax	+resample	+resample+cost
Attribute	68.07	69.43	70.16
Edge	0.72	0.73	0.73
Methods	Triplet	+resample	+resample+cost
Attribute	71.29	71.75	72.43
Edge	0.73	0.73	0.74
Methods	Quintuplet	+resample	+resample+cost
Attribute	81.31	83.39	84.26
Edge	0.76	0.77	0.78

Table 5. Control experiments on class imbalance on MNIST-rot-back-image dataset [26]. The mean per-class accuracy (%) is reported for each case.

Remove (%)	0	20	40
Triplet+resample+cost	76.12	67.18	56.49
Quintuplet	77.62	72.26	65.27
Quintuplet+resample+cost	77.64	75.58	70.13

our full method degrades much more gracefully with increasing imbalance than others. This again validates the efficacy of our quintuplet loss and applied schemes.

6. Conclusion

Class imbalance is common in many vision tasks. Contemporary deep representation learning methods typically adopt class re-sampling or cost-sensitive learning. Through extensive experiments, we have validated their usefulness and further demonstrated that the proposed quintuplet sampling with triple-header loss works remarkably well for imbalanced learning. Our method has been shown superior to the triplet loss, which is commonly adopted for large margin learning but does not enforce the inter-cluster margins in quintuplets. Generalization to higher-order relationships beyond explicit clusters is a future direction to explore.

Acknowledgment. This work is partially supported by SenseTime Group Limited and the Hong Kong Innovation and Technology Support Programme.

References

- [1] P. Arbelaez, M. Maire, C. Fowlkes, and J. Malik. Contour detection and hierarchical image segmentation. *TPAMI*, 33(5):898–916, 2011.
- [2] P. Arbeláez, J. Pont-Tuset, J. Barron, F. Marques, and J. Malik. Multiscale combinatorial grouping. In *CVPR*, 2014.
- [3] T. Berg and P. N. Belhumeur. POOF: Part-Based One-vs-One Features for fine-grained categorization, face verification, and attribute estimation. In *CVPR*, 2013.
- [4] G. Bertasius, J. Shi, and L. Torresani. Deepedge: A multi-scale bifurcated deep network for top-down contour detection. In *CVPR*, 2015.
- [5] G. Bertasius, J. Shi, and L. Torresani. High-for-low and low-for-high: Efficient boundary detection from deep object features and its applications to high-level vision. In *ICCV*, 2015.
- [6] Y.-L. Boureau, N. Le Roux, F. Bach, J. Ponce, and Y. LeCun. Ask the locals: Multi-way local pooling for image recognition. In *ICCV*, 2011.
- [7] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer. Smote: synthetic minority over-sampling technique. *JAIR*, 16(1):321–357, 2002.
- [8] G. Chechik, U. Shalit, V. Sharma, and S. Bengio. An online algorithm for large scale image similarity learning. In *NIPS*, 2009.
- [9] C. Chen, A. Liaw, and L. Breiman. Using random forest to learn imbalanced data. Technical report, University of California, Berkeley, 2004.
- [10] P. Dollár and C. L. Zitnick. Fast edge detection using structured forests. *TPAMI*, 2015.
- [11] C. Drummond and R. C. Holte. C4.5, class imbalance, and cost sensitivity: Why under-sampling beats over-sampling. In *ICMLW*, 2003.
- [12] D. Erhan, Y. Bengio, A. Courville, P.-A. Manzagol, P. Vincent, and S. Bengio. Why does unsupervised pre-training help deep learning? *JMLR*, 11:625–660, 2010.
- [13] Y. Ganin and V. S. Lempitsky. N4-fields: Neural network nearest neighbor fields for image transforms. In *ACCV*, 2014.
- [14] A. Globerson and S. T. Roweis. Metric learning by collapsing classes. In *NIPS*, 2006.
- [15] R. Hadsell, S. Chopra, and Y. LeCun. Dimensionality reduction by learning an invariant mapping. In *CVPR*, 2006.
- [16] S. Hallman and C. Fowlkes. Oriented edge forests for boundary detection. In *CVPR*, 2015.
- [17] H. Han, W.-Y. Wang, and B.-H. Mao. Borderline-SMOTE: A new over-sampling method in imbalanced data sets learning. In *ICIC*, 2005.
- [18] H. He and E. A. Garcia. Learning from imbalanced data. *TKDE*, 21(9):1263–1284, 2009.
- [19] C. Huang, C. C. Loy, and X. Tang. Discriminative sparse neighbor approximation for imbalanced learning. *arXiv preprint, arXiv:1602.01197*, 2016.
- [20] P. Isola, D. Zoran, D. Krishnan, and E. H. Adelson. Crisp boundary detection using pointwise mutual information. In *ECCV*, 2014.
- [21] P. Jeatrakul, K. Wong, and C. Fung. Classification of imbalanced data by combining the complementary neural network and SMOTE algorithm. In *ICONIP*, 2010.
- [22] S. H. Khan, M. Bennamoun, F. Sohel, and R. Togneri. Cost sensitive learning of deep feature representations from imbalanced data. *arXiv preprint, arXiv:1508.03422v1*, 2015.
- [23] J. J. Kivinen, C. K. I. Williams, and N. Heess. Visual boundary prediction: A deep neural prediction network and quality dissection. In *AISTATS*, 2014.
- [24] N. Kumar, A. C. Berg, P. N. Belhumeur, and S. K. Nayar. Attribute and simile classifiers for face verification. In *ICCV*, 2009.
- [25] N. Kumar, A. C. Berg, P. N. Belhumeur, and S. K. Nayar. Describable visual attributes for face verification and image search. *TPAMI*, 33(10):1962–1977, 2011.
- [26] H. Larochelle, D. Erhan, A. Courville, J. Bergstra, and Y. Bengio. An empirical evaluation of deep architectures on problems with many factors of variation. In *ICML*, 2007.
- [27] J. Lim, C. L. Zitnick, and P. Dollár. Sketch tokens: A learned mid-level representation for contour and object detection. In *CVPR*, 2013.
- [28] W. Liu and S. Chawla. Class confidence weighted kNN algorithms for imbalanced data sets. In *PAKDD*, 2011.
- [29] Z. Liu, P. Luo, X. Wang, and X. Tang. Deep learning face attributes in the wild. In *ICCV*, 2015.
- [30] T. Maciejewski and J. Stefanowski. Local neighbourhood extension of SMOTE for mining imbalanced data. In *CIDM*, 2011.
- [31] R. Min, D. A. Stanley, Z. Yuan, A. Bonner, and Z. Zhang. A deep non-linear feature mapping for large-margin kNN classification. In *ICDM*, 2009.
- [32] M. Oquab, L. Bottou, I. Laptev, and J. Sivic. Learning and transferring mid-level image representations using convolutional neural networks. In *CVPR*, 2014.
- [33] X. Ren and L. Bo. Discriminatively Trained Sparse Code Gradients for Contour Detection. In *NIPS*, 2012.
- [34] F. Schroff, D. Kalenichenko, and J. Philbin. Facenet: A unified embedding for face recognition and clustering. In *CVPR*, 2015.
- [35] W. Shen, X. Wang, Y. Wang, X. Bai, and Z. Zhang. Deep-contour: A deep convolutional feature learned by positive-sharing loss for contour detection. In *CVPR*, 2015.
- [36] C. Silpa-Anan and R. Hartley. Optimised KD-trees for fast image descriptor matching. In *CVPR*, 2008.
- [37] K. Simonyan and A. Zisserman. Very deep convolutional networks for large-scale image recognition. In *ICLR*, 2015.
- [38] N. Srivastava and R. R. Salakhutdinov. Discriminative transfer learning with tree-based priors. In *NIPS*, 2013.
- [39] Y. Sun, Y. Chen, X. Wang, and X. Tang. Deep learning face representation by joint identification-verification. In *NIPS*, 2014.
- [40] Y. Tang, Y.-Q. Zhang, N. Chawla, and S. Krasser. SVMs modeling for highly imbalanced classification. *TSMC*, 39(1):281–288, 2009.
- [41] K. M. Ting. A comparative study of cost-sensitive boosting algorithms. In *ICML*, 2000.
- [42] J. Wang, Y. Song, T. Leung, C. Rosenberg, J. Wang, J. Philbin, B. Chen, and Y. Wu. Learning fine-grained image similarity with deep ranking. In *CVPR*, 2014.

- [43] J. Wang, J. Yang, K. Yu, F. Lv, T. Huang, and Y. Gong. Locality-constrained linear coding for image classification. In CVPR, 2010.
- [44] K. Q. Weinberger and L. K. Saul. Distance metric learning for large margin nearest neighbor classification. JMLR, 10:207–244, 2009.
- [45] S. Xie and Z. Tu. Holistically-nested edge detection. In ICCV, 2015.
- [46] B. Zadrozny, J. Langford, and N. Abe. Cost-sensitive learning by cost-proportionate example weighting. In ICDM, 2003.
- [47] N. Zhang, M. Paluri, M. Ranzato, T. Darrell, and L. Bourdev. PANDA: Pose aligned networks for deep attribute modeling. In CVPR, 2014.
- [48] Z.-H. Zhou and X.-Y. Liu. Training cost-sensitive neural networks with methods addressing the class imbalance problem. TKDE, 18(1):63–77, 2006.